

CPSC 532P / LING 530A: Deep Learning for Natural Language Processing (DL-NLP)

Muhammad Abdul-Mageed

`muhammad.mageed@ubc.ca`

Natural Language Processing Lab

The University of British Columbia

Table of Contents

1 Neural Machine Translation

2 Attention

3 Visual Attention

Neural Machine Translation

Machine Translation

- **Phrase-Based Machine Translation (PBMT)**: Many small sub-components that are tuned separately (e.g., Koehn et al. [2003])

Machine Translation

- **Neural Machine Translation (NMT)**: A large neural network that reads a sentence and outputs a translation
- **Encoder-decoder approach** (e.g. Sutskever et al. (2014)):
 - **Encoder** encodes a source sentence into a *fixed-length* vector
 - **Decoder** outputs a translation from the encoded vector
- **Encoder–decoder system** **jointly trained** to maximize the probability of a correct translation given a source sentence

Example Setup

Jaguar F-Type, du rallye dans l'air!

Par Jules Humbert | Publié le 16/11/2018 à 11:43



NOUVEAUTÉ - Jaguar rend hommage aux 70 ans de la naissance de ses modèles XK à travers la réalisation de deux F-Type roadster préparées pour le rallye.

Figure: From Le Figaro.

MT as Seq2Seq Learning: English-French



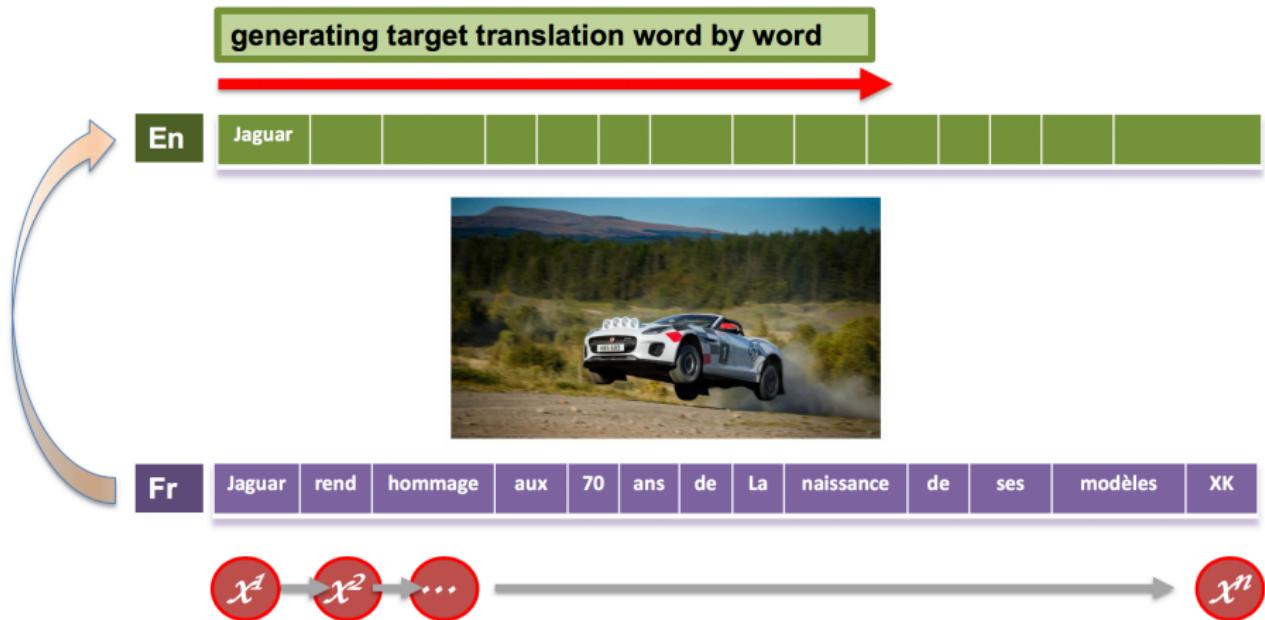
En Jaguar | pays | homage | to | the | 70 | years | of | the | birth | of | its | XK | models

Fr Jaguar | rend | hommage | aux | 70 | ans | de | La | naissance | de | ses | modèles | XK

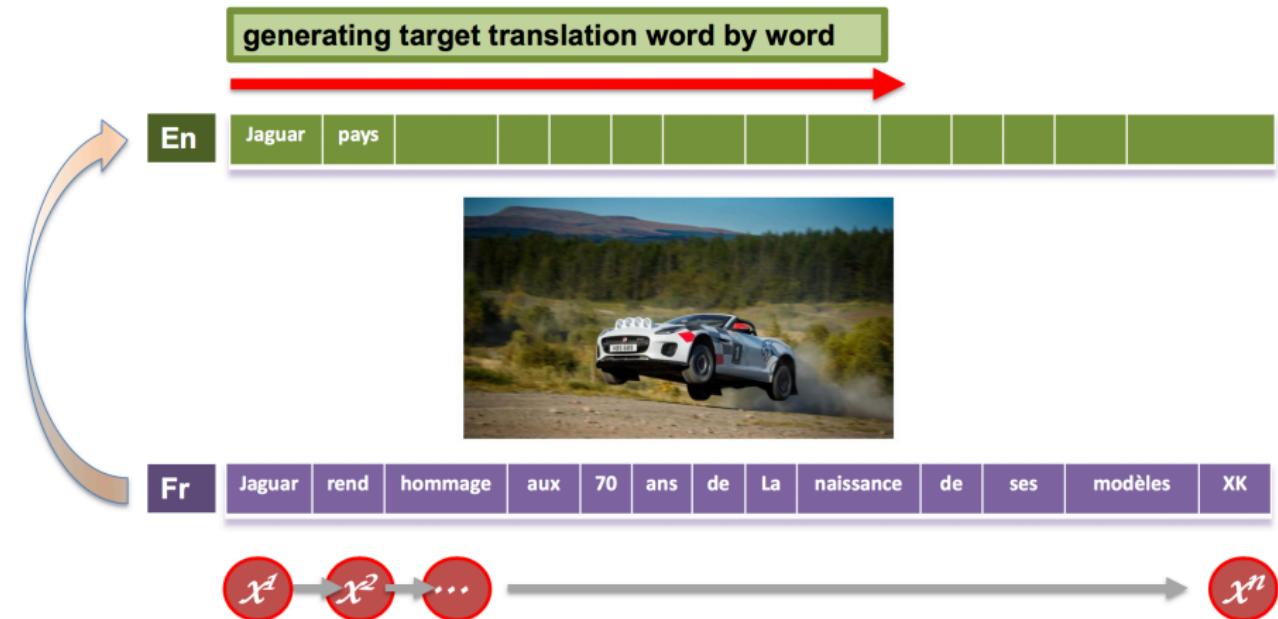
A blue curved arrow points from the English input sequence to the French output sequence.

Figure: Translation by Google, November, 20, 2018.

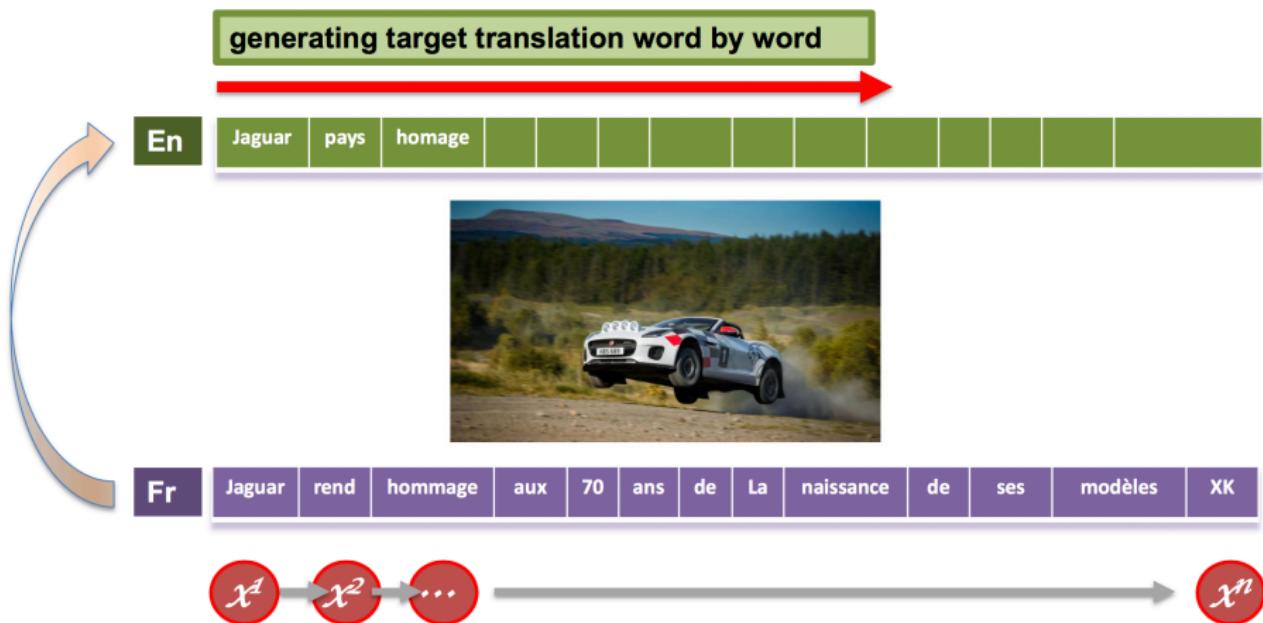
Generating Target (En) One Word at a Time: w_1



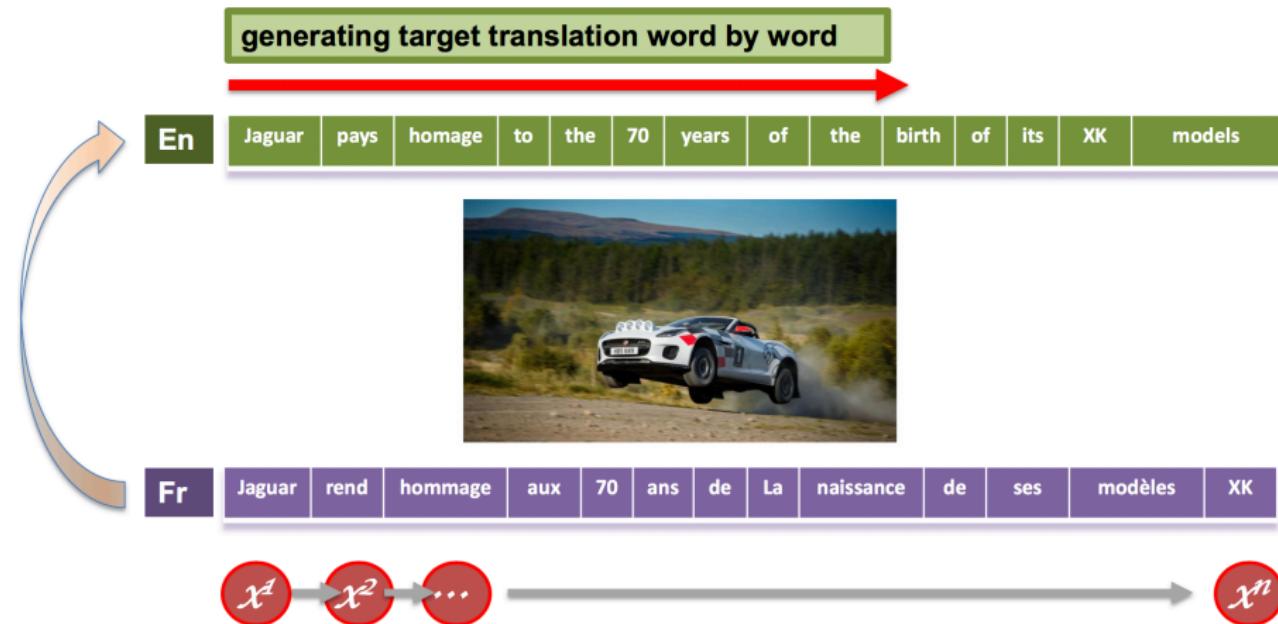
Generating Target (En) One Word at a Time: $w_1 w_2$



Generating Target (En) One Word at a Time: $w_1 w_2 w_3$



Generating Target (En) One Word at a Time: $w_1 w_2 \dots w_n$



1: Encoder

Encoder reads the input sentence \mathbf{x} in to a vector \mathbf{c}

$$\mathbf{x} = (x_1, \dots, x_T)$$

Common to use an **RNN**:

$$h_t = f(x_t, h_{t-1})$$

and

$$\mathbf{c} = q(h_1, \dots, h_{T_x})$$

- where h_t is a **hidden state** at time step t and \mathbf{c} **vector of hidden states**, and f and q are **nonlinear functions** (e.g., LSTM)

Decoder Details

Decoder

- Decoder trained to predict next word y_t , given (1) the context vector \mathbf{c} and (2) all the previously predicted words $\{y_1, \dots, y_{t-1}\}$
- i.e., The decoder defines a probability over the translation \mathbf{y} by decomposing the joint probability into the ordered conditionals:

2: Decoder

$$p(\mathbf{y}) = \prod_{t=1}^T p(y_t | \{y_1, \dots, y_{t-1}\}, \mathbf{c})$$

where (\mathbf{y} is gold word):

$$\mathbf{y} = (y_1, \dots, y_{T_y})$$

3: Decoder *Contd.*

With an RNN, each conditional probability is modelled as:

$$p(y_t | \{y_1, \dots, y_{t-1}\}, c) = g(y_{t-1}, s_t, c)$$

where ***g*** is a nonlinear function that outputs the probability of y_t and ***s_t*** is the hidden state of the RNN.

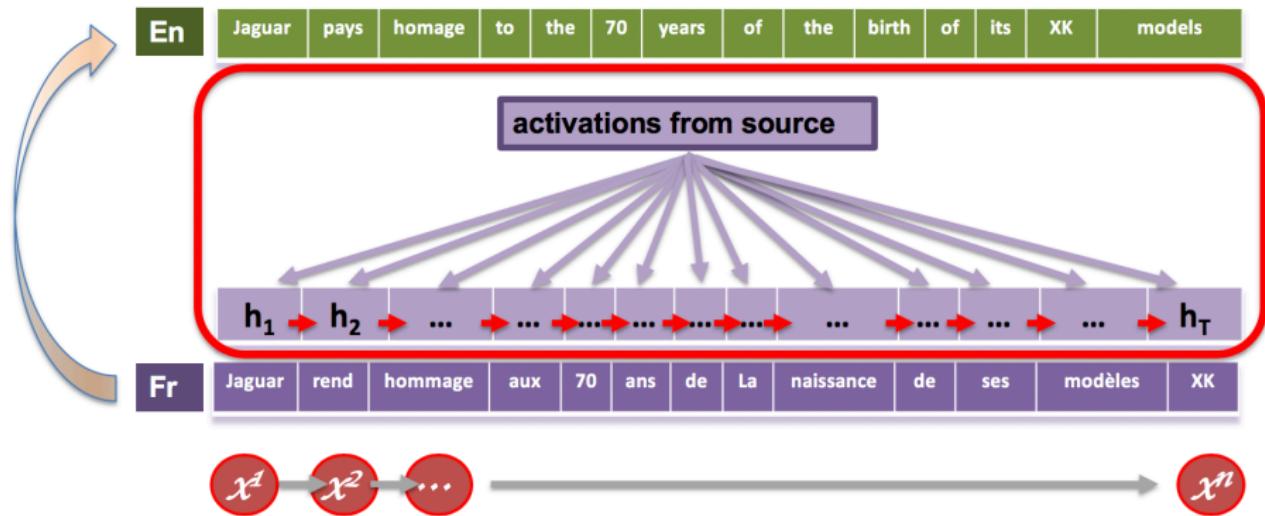
Compressing Long Sentences

- **Issue:** Difficult to compress all the necessary information of a source sentence into a **fixed-length vector**, especially for **long sentences** (Cho et al., 2014)

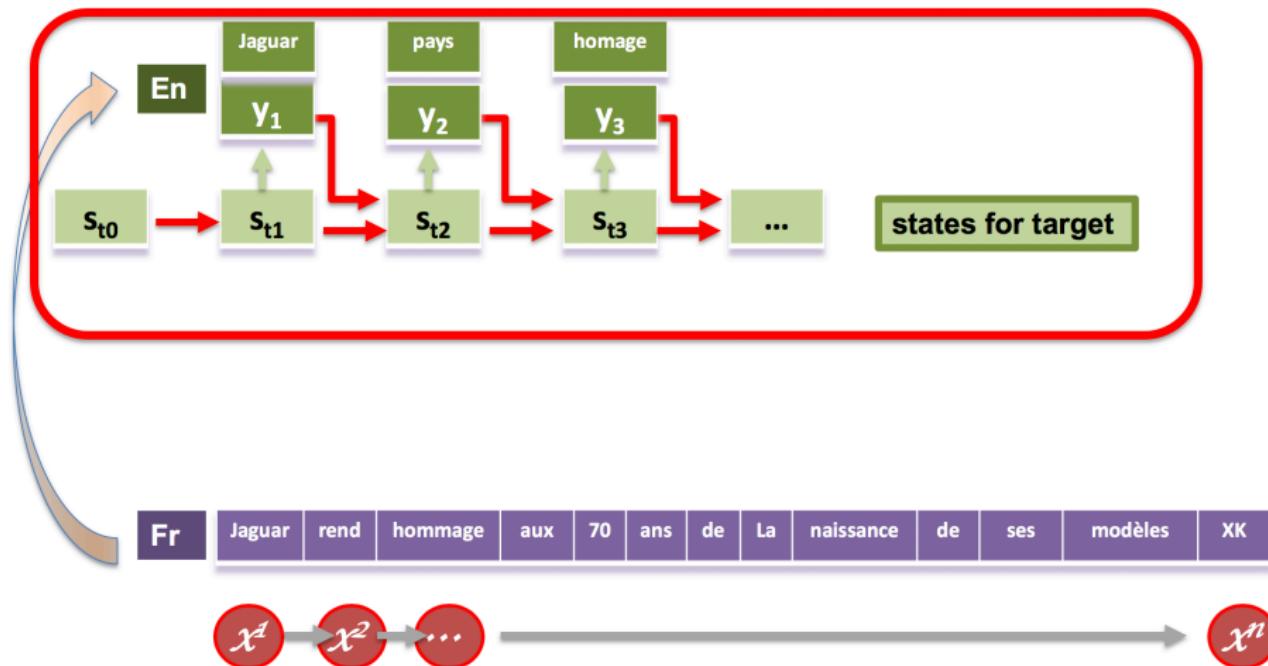
Bahdanau et al. (2015): Learn to ‘align’ and translate jointly

- Generate a word in a translation,
- (**Soft-**)search for a set of positions in a source sentence where the most relevant information is concentrated
- Predict a target word based on the context vectors associated with (1) these source positions and (2) all the previous generated target words
- Copes better with long sentences

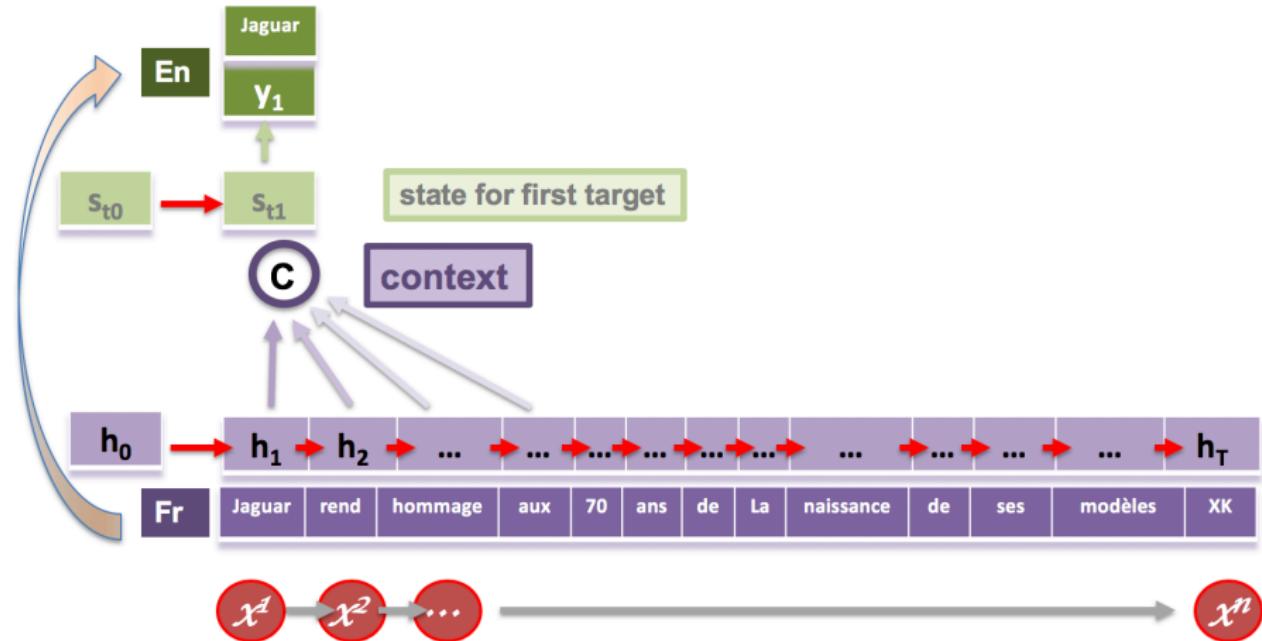
Source Language Activations



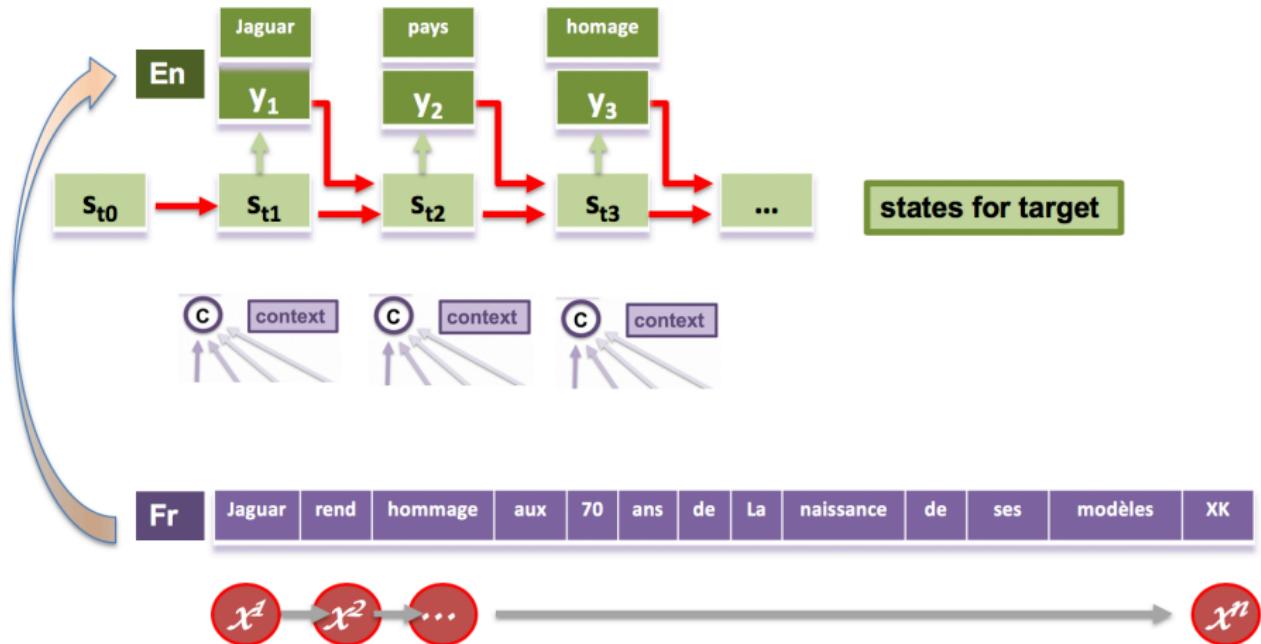
Target (Language) States



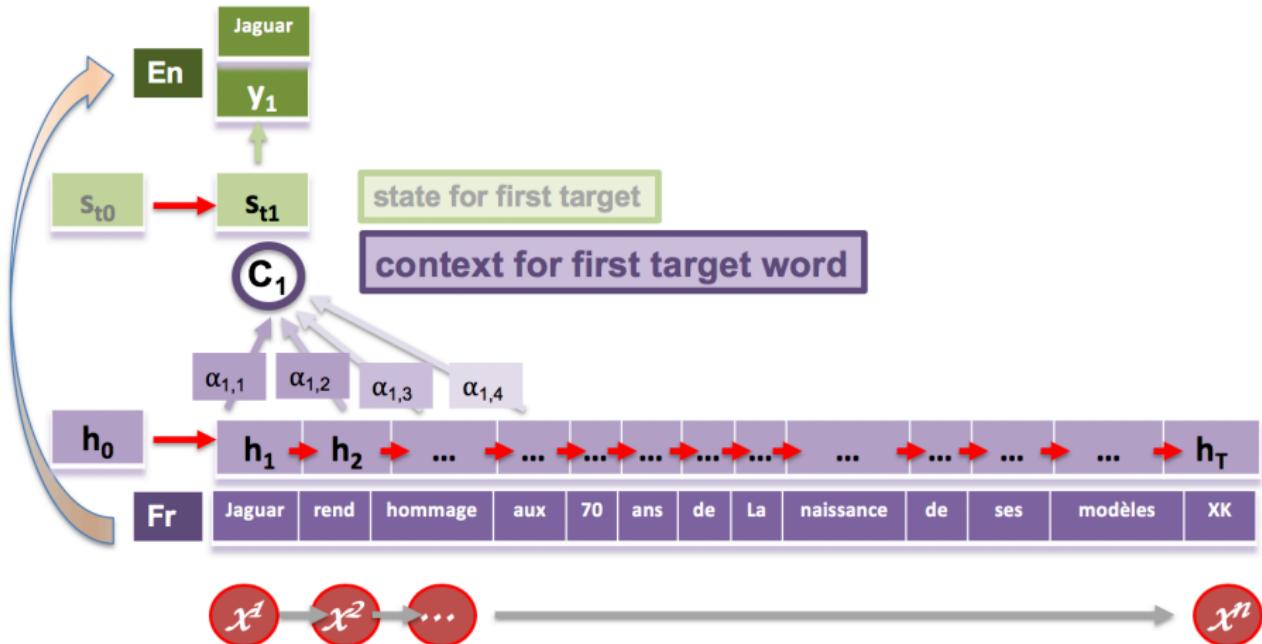
Context



Context Per Every Target State



Context For State One



What is in Context C1?

4: Context C1

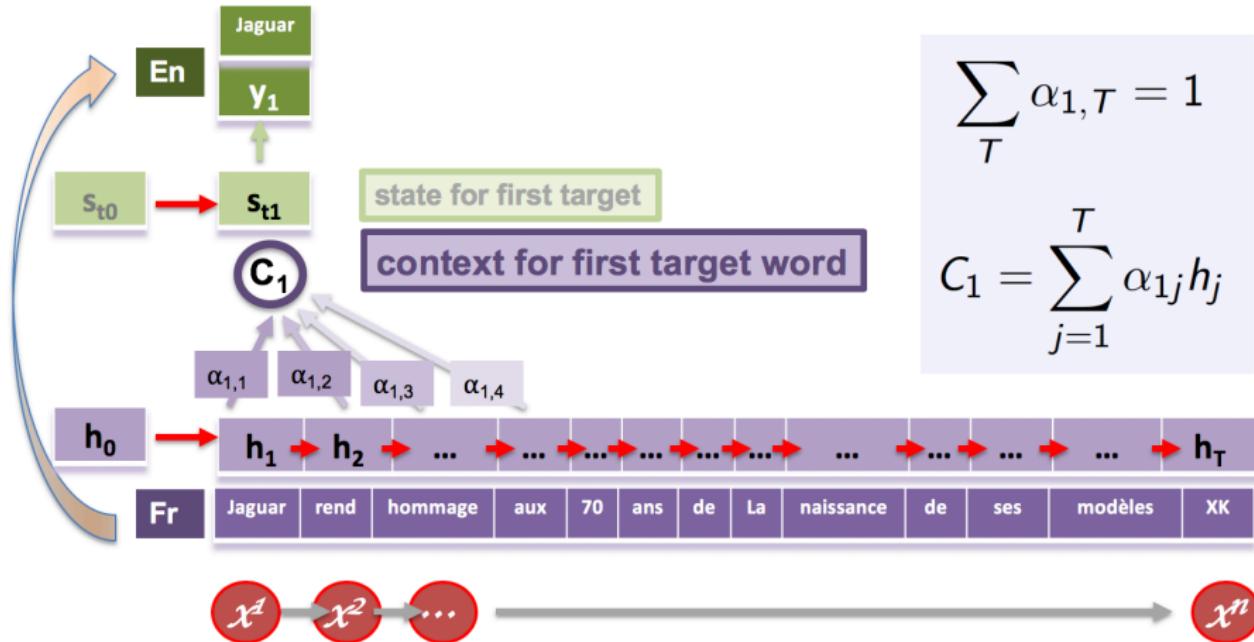
$$\sum_{\tau} \alpha_{1,\tau} = 1$$

$$C_1 = \sum_{j=1}^T \alpha_{1j} h_j$$

Context Ingredients

- α : attention weights
- h : activation from source

Context For State One Enhanced



What is in Context C2?

5: Context C2

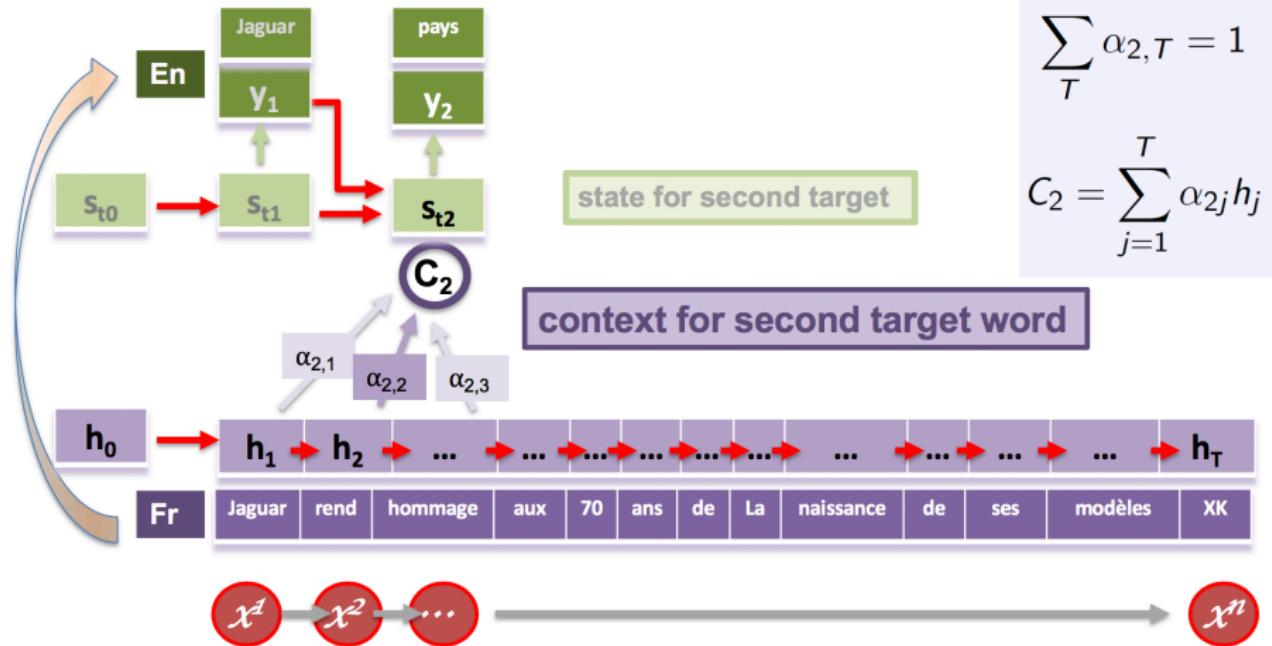
$$\sum_T \alpha_{2,T} = 1$$

$$C_2 = \sum_{j=1}^T \alpha_{2j} h_j$$

Context Ingredients

- α : attention weights
- h : activation from source

Context For State Two



What is in Context C3?

6: Context C3

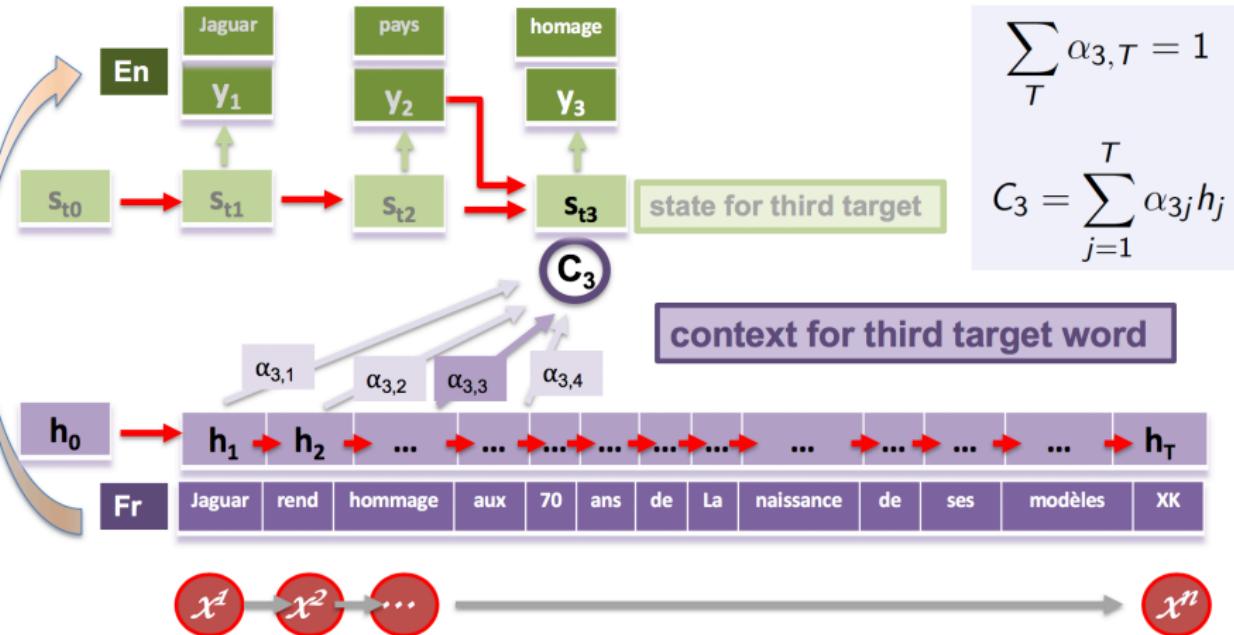
$$\sum_T \alpha_{3,T} = 1$$

$$C_3 = \sum_{j=1}^T \alpha_{3j} h_j$$

Context Ingredients

- α : attention weights
- h : activation from source

Context For State Three



7: New Decoder

Recall, original decoder:

$$p(\mathbf{y}) = \prod_{t=1}^T p(y_t | \{y_1, \dots, y_{t-1}\}, \mathbf{c})$$

For new decoder, each conditional probability defined as:

$$p(y_i | y_1, \dots, y_{i-1}, \mathbf{x}) = g(y_{i-1}, \mathbf{s}_i, \mathbf{c}_i)$$

where \mathbf{s}_i is an RNN hidden state for time i , computed by:

$$\mathbf{s}_i = f(\mathbf{c}_{i-1}, y_{i-1}, \mathbf{c}_i)$$

What is in Context C?

Context Ingredients

- The context vector c_i depends on a sequence of annotations (h_1, \dots, h_{T_x}) to which the encoder maps the input sequence.

8: Context C

$$c_i = \sum_{j=1}^{T_x} \alpha_{ij} h_j$$

Context Ingredients

- α : attention weights
- h : activation from source

Weight α_{ij}

9: Weight α_{ij} for each annotation h_j

$$\alpha_{ij} = \frac{\exp(e_{ij})}{\sum_{k=1}^{T_x} \exp(e_{ik})}$$

where

$$e_{ij} = a(s_{i-1}, h_j)$$

is an **alignment model** based on the RNN hidden state s_{i-1} (just before emitting y_i) and the j -th annotation h_j of the input sentence

Alignment Model Job

- **alignment model a** scores how well the **inputs around position j** and the **output at position i** match.

Computing Soft Alignment for Model α

Alignment Model

- The **alignment model** directly computes a soft alignment, which allows the gradient of the cost function to be backpropagated through
- This gradient can be used to train the alignment model as well as the whole translation model jointly

Role of Attention

- "By letting the decoder have an attention mechanism, we relieve the encoder from the burden of having to encode all information in the source sentence into a fixed-length vector. With this new approach the information can be spread throughout the sequence of annotations, which can be selectively retrieved by the decoder accordingly"
(Bahdanau et al., 2015)

Bahdanau et al. (2015) Attention Model Visualization

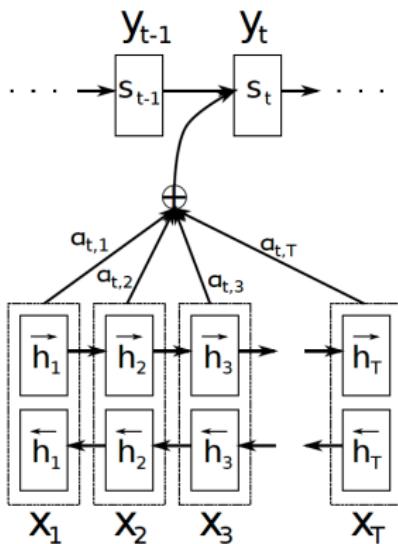


Figure: Bahdanau et al. (2015) attention model trying to generate the t -th target word y_t given a source sentence (x_1, x_2, \dots, x_T)

Bahdanau et al. (2015) Weight Visualization

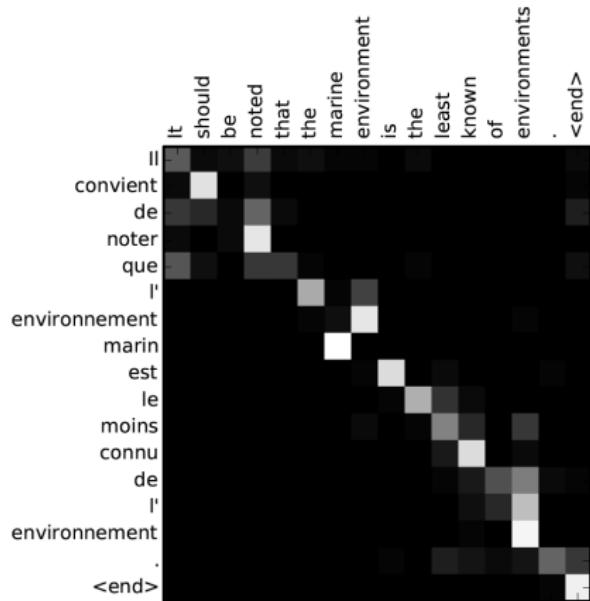
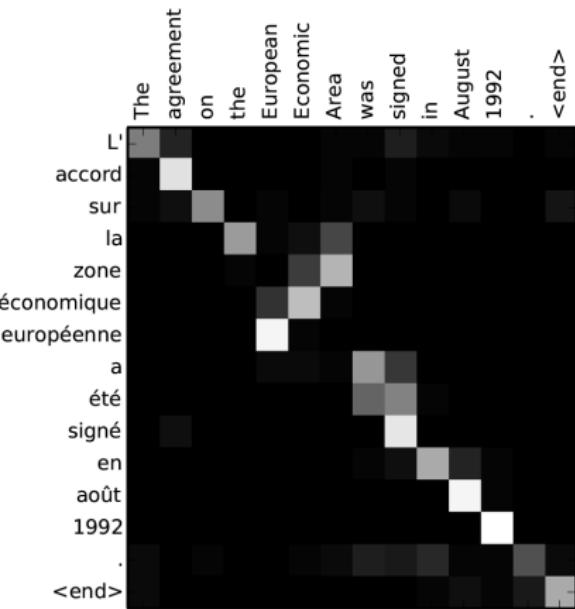


Figure: Each pixel shows the weight ij of the annotation of the j -th source word for the i -th target word (Bahdanau et al., 2015).

Image Caption Generation (With Attention)

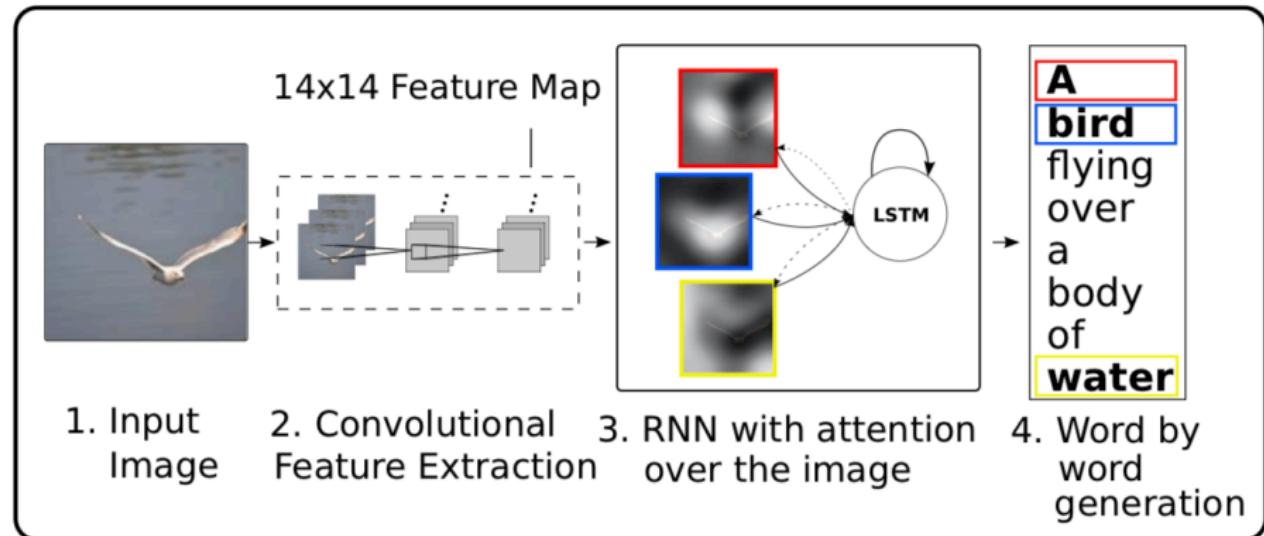


Figure: Xu et al. (2016) Show, Attend, and Tell A model that learns a words/image alignment.

Visual Attention



A woman is throwing a frisbee in a park.

A dog is standing on a hardwood floor.

A stop sign is on a road with a mountain in the background.



A little girl sitting on a bed with a teddy bear.



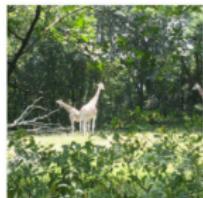
A group of people sitting on a boat in the water.



A giraffe standing in a forest with trees in the background.

Figure: Xu et al. (2016): Show, Attend, and Tell: Neural Image Caption Generation with Visual Attention

Visual Attention: Model Errors



A large white bird standing in a forest.



A woman holding a clock in her hand.



A man wearing a hat and a hat on a skateboard.



A person is standing on a beach with a surfboard.



A woman is sitting at a table with a large pizza.



A man is talking on his cell phone while another man watches.

Figure: Errors from Xu et al. (2016)

Coming Up: Transformer

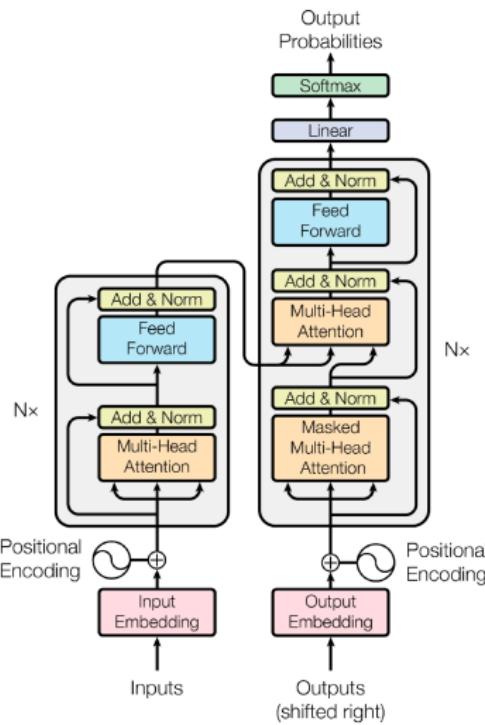


Figure: Transformer. Self-attention is key (Vaswani et al., 2017)